



Monitoring and Optimizing Spark Job

Dano Lee

Catalyst Optimizer

Catalyst Optimizer

- ❖ This is a rule based engine.
- ❖ It takes the logical plan which we expressed in a Dataframe API and then re-writes it into an Optimized Physical Plan. Physical Plan is developed before the query is executed.
- ❖ The catalyst optimizer is at the core of the Spark SQL's power and speed. It automatically finds the most efficient plan for applying your transformations and actions.

Benefits of Catalyst Optimizer

- ❖ **Optimize Query Plans**

- ❖ It analyzes and transforms the logical execution plans of SQL queries to generate optimized physical plans that can be executed efficiently.

- ❖ **Enable Advanced Features**

- ❖ It supports various advanced features such as predicate pushdown, join reordering, constant folding, and more.

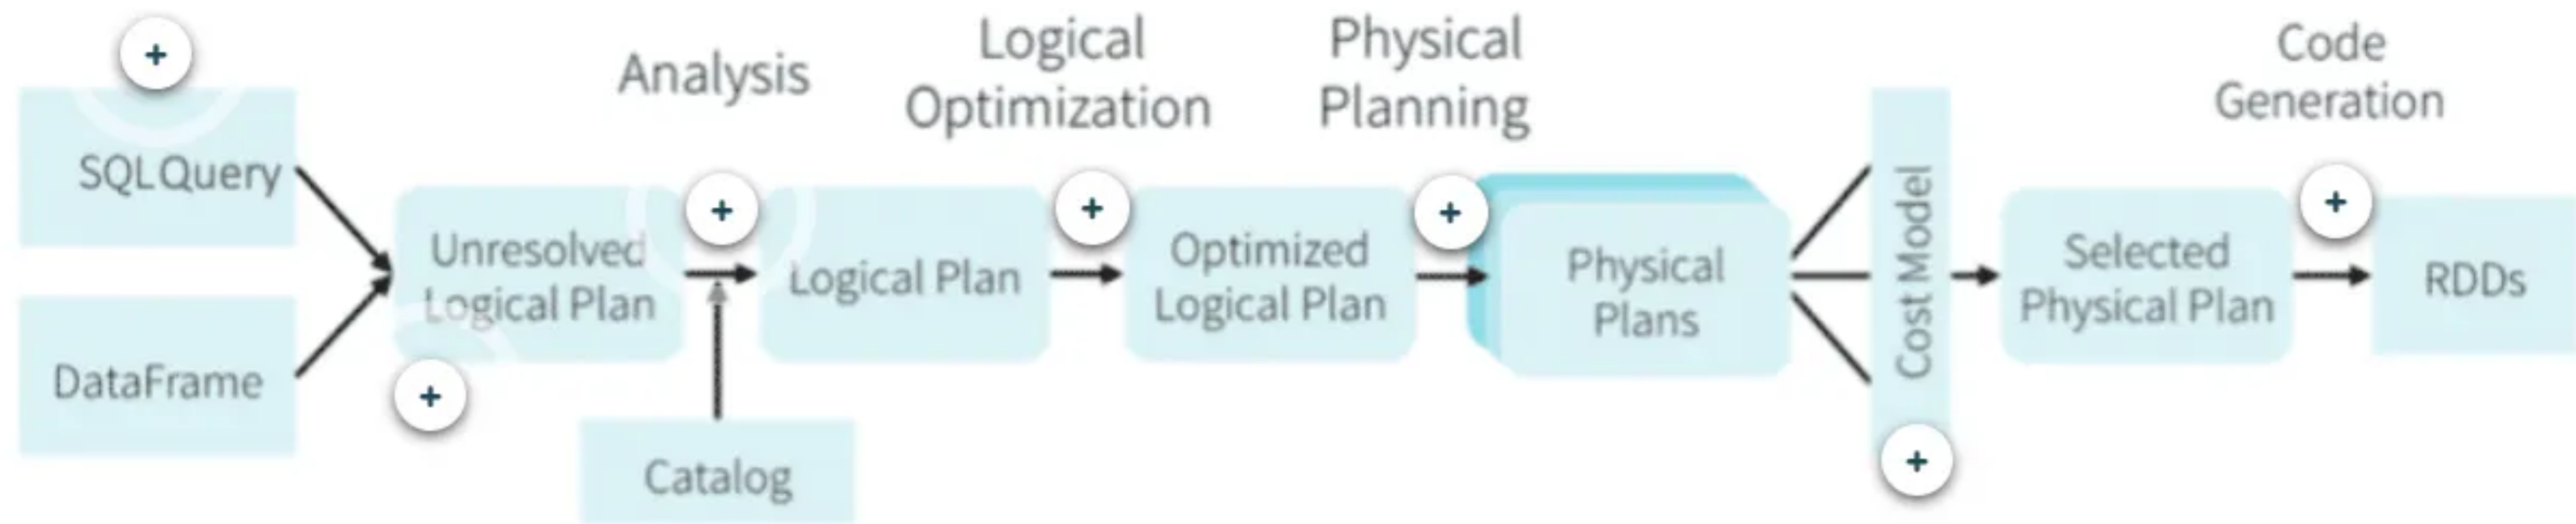
- ❖ **Code Generation**

- ❖ Catalyst also works closely with Tungsten, Spark's execution engine, to generate optimized code for the execution of these plans.

Catalyst Optimizer

- ❖ Let's suppose you write the code in Java, Scala, Python and R in the primary language supported by Apache Spark. When we use this language we use Dataframe or Dataset API. **These are declarative API** that is telling Apache Spark **what we intend to do and NOT how to do.**
- ❖ As Apache Spark SQL is a declarative API is another way of expressing what we want the Apache Spark to do for us and NOT how to do.

Catalyst Optimizer



How Catalyst Optimizer Works Internally (1)

- ❖ **Parsing:**

- ❖ The SQL query or DataFrame operation is parsed into an Abstract Syntax Tree (AST).

- ❖ **Logical Plan Generation:**

- ❖ The AST is converted into a Logical Plan. The logical plan is a tree of relational operators (e.g., Filter, Project, Join) that describes the computation required by the query.

- ❖ **Logical Plan Optimization:**

- ❖ The logical plan undergoes several optimization phases where Catalyst applies various optimization rules. These include:
 - ❖ **Constant Folding:** Simplifies expressions by evaluating constant expressions at compile time.
 - ❖ **Predicate Pushdown:** Moves filters as close to the data source as possible, reducing the amount of data that needs to be processed.
 - ❖ **Join Reordering:** Reorders joins to optimize the execution by minimizing the size of intermediate results.
 - ❖ **Projection Pruning:** Removes unnecessary columns from the query plan, reducing the amount of data read and processed.

How Catalyst Optimizer Works Internally (2)

- ❖ **Physical Plan Generation:**

- ❖ The optimized logical plan is then converted into one or more Physical Plans. Each physical plan corresponds to a specific way to execute the query, considering different strategies like different join algorithms (e.g., broadcast join, sort-merge join).

- ❖ **Physical Plan Optimization:**

- ❖ Catalyst compares different physical plans and chooses the most efficient one based on a cost model, which estimates the computational resources needed (e.g., CPU, I/O).

- ❖ **Code Generation:**

- ❖ Finally, the selected physical plan is compiled into optimized Java bytecode using **Tungsten's** code generation capabilities, which is then executed by Spark's runtime engine.

Dynamic Partition Pruning

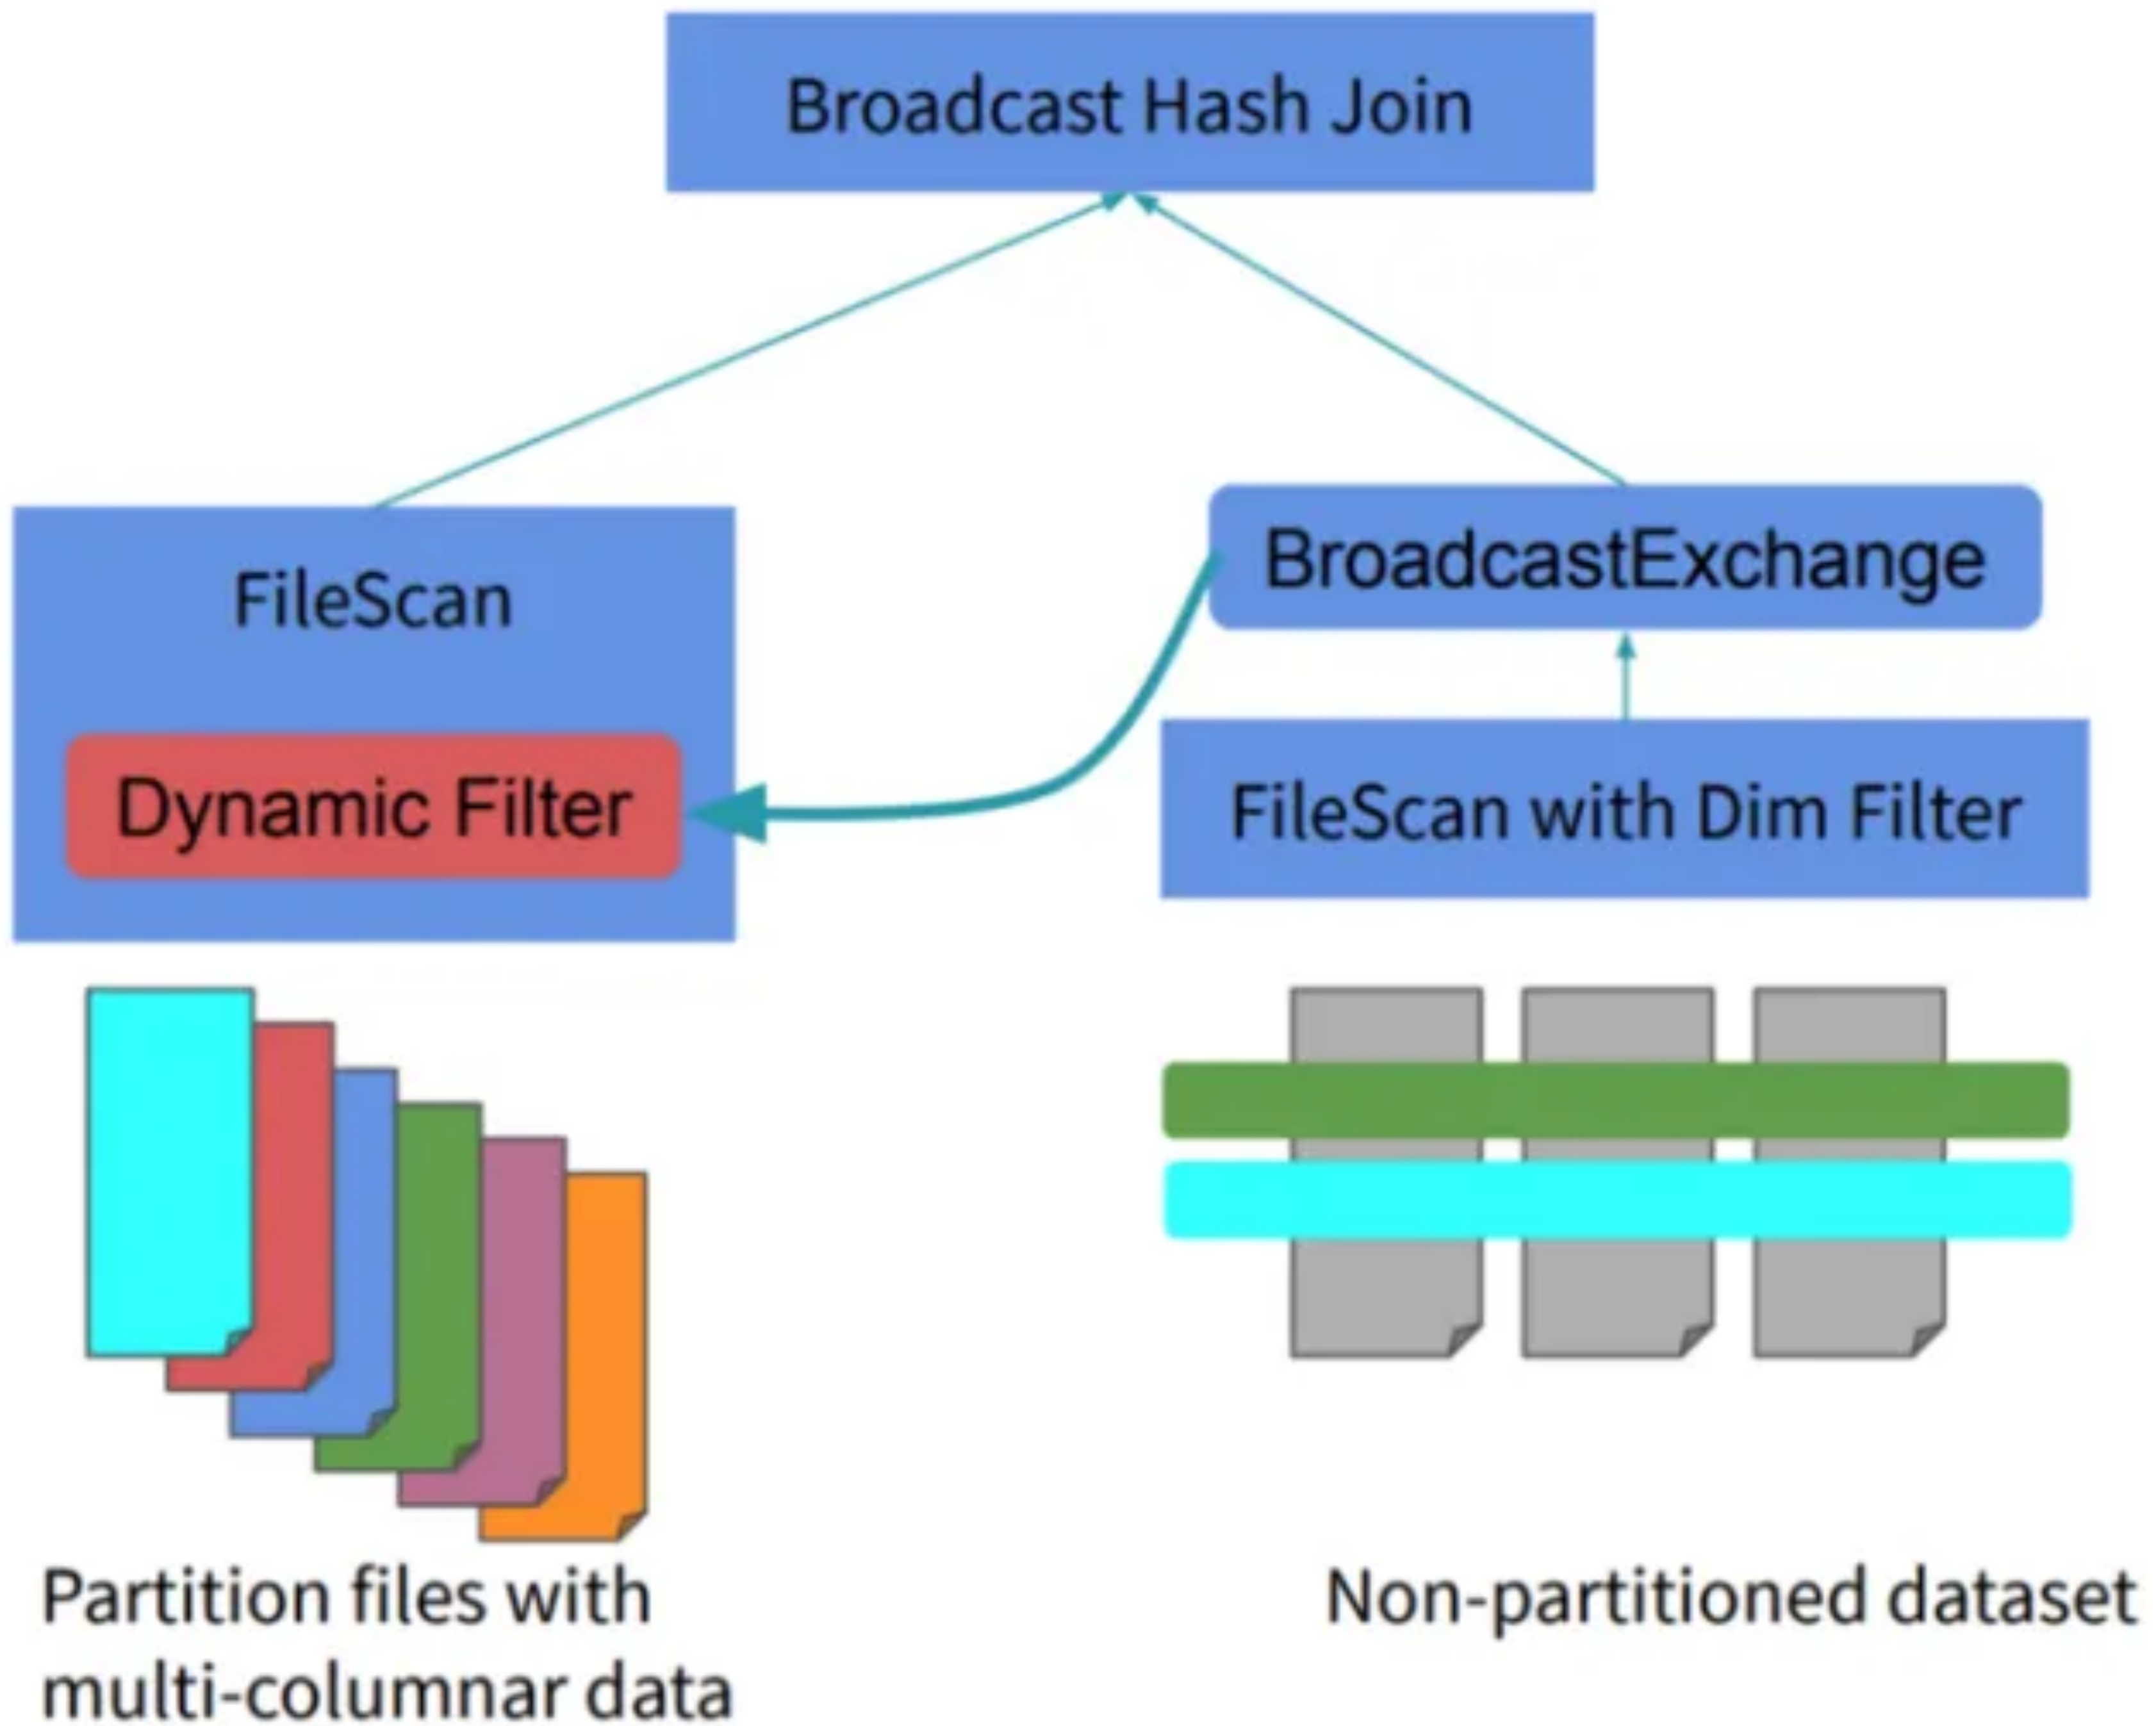
- ❖ **Dynamic Partition Pruning = Predicate Push Down + Broadcast Hash Join**
- ❖ **`spark.sql.optimizer.dynamicPartitionPruning.enabled=true`**

Example of DPP

- ❖ Imagine you have a **fact table with sales records (Sales)** and a **dimension table with dates (Date)**.
- ❖ If you **filter the Date table to only include January 2024 dates**, Spark will use DPP to **only scan partitions in the Sales table that contain records from January 2024**, instead of scanning the entire Sales table.

Dynamic Partition Pruning

- ❖ The smaller(**dimension**) table is queried and filtered. **A hash table is built as part of the filter query.**
- ❖ Spark uses the result of this query (and hash table) to create a broadcast variable
- ❖ Then, it will broadcast the filter to each executor
- ❖ At runtime, Spark's physical plan is changed so that the dynamic filter is applied to the bigger(**fact**) table. This dynamic filter is created as an internal subquery built from the filter applied to the smaller table.
- ❖ The dynamic filter is essentially an internal subquery that Spark creates based on the filter applied to the smaller table. When Spark scans the larger table, it applies this dynamic filter, **skipping partitions that do not meet the criteria**, thereby significantly reducing the amount of data that needs to be processed.



Constraints of Dynamic Partition Pruning

- ❖ Tables that need to be pruned, **must be partitioned with any one of the join key columns.**
- ❖ It works only with Equi-join (joins with the '=' condition).
- ❖ DPP won't apply to the correlated subqueries
- ❖ DPP is useful for queries that follow the Star-schema architectural model.

Spark Web UI

Spark Web UI

- ❖ Spark Driver Web UI
 - ❖ The Spark Driver by default exposes the Spark Web UI using **information from the live application**.
 - ❖ This is the most comprehensive interface, as you can **see live data of your applications in terms of disk and memory utilization**.
 - ❖ You can easily locate and connect to this interface by **connecting on the YARN Resource Manager and then opening the Application Master URL available for the running Spark application**.

Spark History Server

- ❖ This is a service that Spark provides to review completed or running applications on the cluster.
- ❖ It's important to note that this interface **uses the Spark Event Logs** to populate the information that are available only at runtime (e.g. caching). It may miss out on some details that were only accessible while the application was actively running.

Spark History Server

- ❖ `./sbin/start-history-server.sh`
- ❖ The spark jobs themselves must be configured to log events, and to log them to the same shared, writable directory. For example, if the server was configured with a log directory of `hdfs://namenode/shared/spark-logs`, then the client-side options would be:
 - ❖ `spark.eventLog.enabled true`
 - ❖ `spark.eventLog.dir hdfs://namenode/shared/spark-logs`

Retrieve Spark Event Logs

- ❖ When using Amazon EMR, the Spark Event logs are enabled by default and are automatically stored on the HDFS of the cluster where the job was running under the HDFS path `/var/log/spark/apps/`



History Server

Event log directory: `hdfs:///var/log/spark/apps`

Last updated: 2023-09-08 23:05:06

Client local time zone: Europe/Rome

Search:


| Version | App ID | App Name | Started | Completed | Duration | Spark User | Last Updated | Event Log |
|--------------|------------------------------------------------|----------|---------------------|---------------------|----------|------------|---------------------|--------------------------|
| 3.4.0-amzn-0 | application_1694206676971_0001 | Spark Pi | 2023-09-08 22:59:48 | 2023-09-08 23:00:19 | 32 s | hadoop | 2023-09-08 23:00:19 | Download |

Showing 1 to 1 of 1 entries

[Show incomplete applications](#)

Applications in Spark Web UI

- ❖ Specifically the **duration** field is a useful field as it provides a metrics for the duration of the application since the Spark driver was launched and terminated, excluding additional submission details that are specific for different deployment models, so it can be useful to compare the Spark runtime across different versions or providers.

 **History Server**
3.4.0-amzn-0

Event log directory: `hdfs:///var/log/spark/apps`
Last updated: 2023-09-08 23:31:51
Client local time zone: Europe/Rome


Search:

| Version | App ID | App Name | Started | Completed | Duration | Spark User | Last Updated | Event Log |
|--------------|------------------------------------------------|------------------------------------------------------------------|---------------------|---------------------|----------|------------|---------------------|--------------------------|
| 3.4.0-amzn-0 | application_1694208236041_0001 | TPCDS Benchmark - c5d.9xlarge - 1GB - maximizeResourceAllocation | 2023-09-08 23:25:54 | 2023-09-08 23:31:45 | 5.8 min | hadoop | 2023-09-08 23:31:45 | Download |

Showing 1 to 1 of 1 entries
[Show incomplete applications](#)

Jobs in Spark Web UI (1)

- ❖ it can give you a good indication of which portions of your code took more time to execute and you can also use it to compare two jobs executed in different environments or with different configurations to spot differences in terms of time processing. In those last cases is useful to sort your Jobs by Duration using the interface.

 3.4.0-amzn-0

Jobs

Stages

Storage

Environment

Executors

SQL / DataFrame

TPCDS Benchmark - c5d.9xlarge - ... application UI

Spark Jobs (?)

User: hadoop
Total Uptime: 5.8 min
Scheduling Mode: FIFO
Completed Jobs: 930

▶ Event Timeline

▼ Completed Jobs (930)

Page: 1 2 3 4 5 6 7 8 9 10 > >>

93 Pages. Jump to 1. Show 10 items in a page. Go

| Job Id (Job Group) ▾ | Description | Submitted | Duration | Stages: Succeeded/Total | Tasks (for all stages): Succeeded/Total |
|----------------------|------------------------------------------------------------------------------------|---------------------|----------|-------------------------|-----------------------------------------|
| 929 | show at SparkBenchmark.scala:99 show at SparkBenchmark.scala:99 | 2023/09/08 21:31:45 | 0.1 s | 1/1 (1 skipped) | 1/1 (2 skipped) |
| 928 | show at SparkBenchmark.scala:99 show at SparkBenchmark.scala:99 | 2023/09/08 21:31:44 | 0.3 s | 1/1 | 2/2 |

Jobs in Spark Web UI (2)

- ❖ Additionally, this page provides additional information on the top left page:
- ❖ **User*** The user who launched the application. In Amazon EMR this typically match the hadoop user unless you're using the Hadoop impersonation.
- ❖ **Total Uptime*** Time since the Spark application started till the completion of the last Job **Scheduling Mode*** The internal scheduler used within the Spark Application to execute the Spark Jobs. By default Spark uses a FIFO (First In First Out) scheduler.*

APACHE Spark 3.4.0-amzn-0

Jobs Stages Storage Environment Exec

Spark Jobs (?)


User: hadoop
Total Uptime: 5.8 min
Scheduling Mode: FIFO
Completed Jobs: 930
▶ Event Timeline
▼ Completed Jobs (930)

Page: 1 2 3 4 5 6 7 8 9 10 > >>

| Job Id (Job Group) ▼ | Description |
|----------------------|------------------------------------------------------------------------------------|
| 929 | show at SparkBenchmark.scala:99 show at SparkBenchmark.scala:99 |
| 928 | show at SparkBenchmark.scala:99 show at SparkBenchmark.scala:99 |
| 927 | csv at SparkBenchmark.scala:98 csv at SparkBenchmark.scala:98 |


Jobs in Spark Web UI (3)

- ❖ Lastly, this page allows you to review the lifecycle of your application, by expanding the **Event Timeline**** section
 - ❖ you can review how the different Spark Jobs were executed during the time, as also Spark Executors launch and termination, that can give you useful information to detect slow-downs due to the lack of resources (e.g. you're using the Spark Dynamic allocation along with a cluster managed scaler and the nodes took too much time to be added to the cluster).



Stages in Spark Web UI (1)

- ❖ As for the Jobs page, you can review also all the Stages that have been processed in your application. This pages can be reached directly from the Web UI, and in this case it will display all the Stages of the application, or you can select a single Spark Job in the Jobs** section if you're only interested to the Stages processed in an individual Spark Job.

 3.4.0-amzn-0

Jobs **Stages** Storage Environment Executors SQL / DataFrame

TPCDS Benchmark - c5d.9xlarge - ... application UI

Stages for All Jobs

Completed Stages: 930
Skipped Stages: 58

▼ Completed Stages (930)

Page: 1 2 3 4 5 6 7 8 9 10 > >>

93 Pages. Jump to 1 . Show 10 items in a page. Go

| Stage Id ▾ | Description | | Submitted | Duration | Tasks: Succeeded/Total | Input | Output | Shuffle Read | Shuffle Write |
|------------|-------------------------------------------------|--------------------------|---------------------|----------|------------------------|---------|---------|--------------|---------------|
| 1387 | show at SparkBenchmark.scala:99 | +details | 2023/09/08 21:31:45 | 0.1 s | <div>1/1</div> | | | 10.7 KiB | |
| 1385 | show at SparkBenchmark.scala:99 | +details | 2023/09/08 21:31:44 | 0.3 s | <div>2/2</div> | 6.7 MiB | | | 10.7 KiB |
| 1384 | csv at SparkBenchmark.scala:98 | +details | 2023/09/08 21:31:44 | 0.3 s | <div>1/1</div> | | 5.1 KiB | 1945.0 B | |

- ❖ If you expand an individual Stage, you'll be redirected on a more detailed page where you can examine aggregated metrics of the Tasks processed in the stage.

Details for Stage 982 (Attempt 0)

- ▶ DAG Visualization

| Metric | Min | 25th percentile | Median | 75th percentile | Max |
|------------------------------|------------------|------------------|----------------|------------------|-----------------|
| Task Deserialization Time | 27.0 ms | 30.0 ms | 31.0 ms | 34.0 ms | 40.0 ms |
| Duration | 0.2 s | 0.2 s | 0.3 s | 0.3 s | 0.7 s |
| GC Time | 0.0 ms | 0.0 ms | 0.0 ms | 34.0 ms | 0.4 s |
| Result Serialization Time | 0.0 ms | 0.0 ms | 0.0 ms | 0.0 ms | 0.0 ms |
| Getting Result Time | 0.0 ms | 0.0 ms | 0.0 ms | 0.0 ms | 0.0 ms |
| Scheduler Delay | 6.0 ms | 14.0 ms | 19.0 ms | 23.0 ms | 0.3 s |
| Peak Execution Memory | 228.2 MiB | 292.3 MiB | 292.3 MiB | 292.4 MiB | 292.7 MiB |
| Input Size / Records | 37.6 KiB / 559 | 40.2 KiB / 856 | 41.2 KiB / 966 | 51 KiB / 2065 | 67.3 KiB / 3977 |
| Shuffle Write Size / Records | 342.8 KiB / 3127 | 460.3 KiB / 4477 | 509 KiB / 5065 | 870.9 KiB / 9863 | 1.4 MiB / 17546 |
| Shuffle Write Time | 2.0 ms | 2.0 ms | 2.0 ms | 3.0 ms | 0.4 s |

Tasks (366)

Search:

| Index | Task ID | Attempt | Status | Locality level | Executor ID | Host | Logs | Launch Time | Duration | GC Time | Scheduler Delay | Task Deserialization Time | Result Serialization Time | Getting Result Time | Peak Execution Memory | Input Size / Records | Shuffle Write Time | Shuffle Write Size / Records | Errors |
|-------|---------|---------|---------|----------------|-------------|--------------------------------------------|------------------|---------------------|----------|---------|-----------------|---------------------------|---------------------------|---------------------|-----------------------|----------------------|--------------------|------------------------------|--------|
| 0 | 95202 | 0 | SUCCESS | RACK_LOCAL | 22 | ip-172-31-1-166.eu-west-1.compute.internal | stderr stderr | 2023-09-08 23:30:09 | 0.4 s | 82.0 ms | 14.0 ms | 31.0 ms | | | 292.7 MiB | 67.3 KiB / 3977 | 2.0 ms | 1.4 MiB / 17546 | |

Details for Stage 982 (Attempt 0)

Resource Profile Id: 0
Total Time Across All Tasks: 1.8 min
Locality Level Summary: Rack local: 366
Input Size / Records: 16.4 MiB / 551264
Shuffle Write Size / Records: 240.5 MiB / 2678088
Associated Job Ids: [653](#)

- ▶ [DAG Visualization](#)
- ▶ [Show Additional Metrics](#)
- ▶ [Event Timeline](#)

Summary Metrics for 366 Completed Tasks

| Metric | Min | 25th percentile |
|------------------------------|------------------|------------------|
| Task Deserialization Time | 27.0 ms | 30.0 ms |
| Duration | 0.2 s | 0.2 s |
| GC Time | 0.0 ms | 0.0 ms |
| Result Serialization Time | 0.0 ms | 0.0 ms |
| Getting Result Time | 0.0 ms | 0.0 ms |
| Scheduler Delay | 6.0 ms | 14.0 ms |
| Peak Execution Memory | 228.2 MiB | 292.3 MiB |
| Input Size / Records | 37.6 KiB / 559 | 40.2 KiB / 856 |
| Shuffle Write Size / Records | 342.8 KiB / 3127 | 460.3 KiB / 4477 |
| Shuffle Write Time | 2.0 ms | 2.0 ms |

▶ Aggregated Metrics by Executor

Tasks (366)

Show


20

 entries

| Index | Task ID | Attempt | Status | Locality level | Executor ID | Host | Logs | Launch Time | Duration | GC Time |
|-------|---------|---------|---------|----------------|-------------|--------------------------------------------|--------------------------------------------------|---------------------|----------|---------|
| 0 | 95202 | 0 | SUCCESS | RACK_LOCAL | 22 | ip-172-31-1-166.eu-west-1.compute.internal | stderr stdout | 2023-09-08 23:30:09 | 0.4 s | 82.0 ms |

Storage in Spark Web UI

- ❖ TheStorage page** contains information about RDD blocks that have been cached or persisted in memory or on the local disks of the cluster. This page will show some details only if you explicitly invoke a persist or cache operation against a Spark Dataframe. More in detail, the page shows for each RDD:

|  3.4.0-amzn-0 | Jobs | Stages | Storage | Environment | Executors | SQL / DataFrame | Spark shell application UI |
|-------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------|-------------------|-----------------|----------------|-----------------|----------------------------|
| Storage | | | | | | | |
| ▼ RDDs | | | | | | | |
| ID | RDD Name | Storage Level | Cached Partitions | Fraction Cached | Size in Memory | Size on Disk | |
| 9 | <div><div>*(1) ColumnarToRow +- FileScan parquet</div><div>[ss_sold_time_sk#0,ss_item_sk#1,ss_customer_sk#2,ss_cdemo_sk#3,ss_hdemo_sk#4,ss_addr_sk#5,ss_store_sk#6,ss_promo_sk#7,ss_ticket_number#8L,ss_quantity#9,ss_wholesale_cost#10,ss_list_price#11,ss_sales_price#12,ss_ext_discount_amt#13,ss_ext_sales_price#14,ss_ext_wholesale_cost#15,ss_ext_list_price#16,ss_ext_tax#17,ss_coupon_amt#18,ss_net_paid#19,ss_net_paid_inc_tax#20,ss_net_profit#21,ss_sold_date_sk#22] Batched: true, DataFilters: [], Format: Parquet, Location: InMemoryFileIndex(1 paths)[s3://ripani.dub/warehouse/tpcds_parquet_1tb/store_sales], PartitionFilters: [], PushedFilters: [], ReadSchema: struct<ss_sold_time_sk:int,ss_item_sk:int,ss_customer_sk:int,ss_cdemo_sk:int,ss_hdemo_sk:int,ss_a...</div></div> | Disk Memory Serialized 2x Replicated | 1275 | 100% | 156.6 GiB | 452.8 GiB | |

RDD Storage Info for *(1) ColumnarToRow +- FileScan parquet [ss_sold_time_sk#0,ss_item_sk#1,ss_customer_sk#2,ss_cdemo_sk#...

Storage Level: Disk Memory Serialized 2x Replicated

Cached Partitions: 1278

Total Partitions: 1278

Memory Size: 156.6 GiB

Disk Size: 459.9 GiB

Data Distribution on 6 Executors

| Host | On Heap Memory Usage | Off Heap Memory Usage | Disk Usage |
|--------------------------------------------------|--------------------------------|-------------------------|------------|
| ip-172-31-3-222.eu-west-1.compute.internal:44907 | 26.2 GiB (336.4 MiB Remaining) | 0.0 B (0.0 B Remaining) | 89.1 GiB |
| ip-172-31-3-244.eu-west-1.compute.internal:43799 | 26.3 GiB (267.1 MiB Remaining) | 0.0 B (0.0 B Remaining) | 44.3 GiB |
| ip-172-31-3-170.eu-west-1.compute.internal:46697 | 25.8 GiB (751.3 MiB Remaining) | 0.0 B (0.0 B Remaining) | 91.4 GiB |
| ip-172-31-3-36.eu-west-1.compute.internal:46023 | 26.1 GiB (476.3 MiB Remaining) | 0.0 B (0.0 B Remaining) | 94.3 GiB |
| ip-172-31-3-39.eu-west-1.compute.internal:43111 | 26.0 GiB (594.8 MiB Remaining) | 0.0 B (0.0 B Remaining) | 49.5 GiB |
| ip-172-31-3-196.eu-west-1.compute.internal:39167 | 26.2 GiB (344.6 MiB Remaining) | 0.0 B (0.0 B Remaining) | 91.4 GiB |

1278 Partitions

| Block Name ▾ | Storage Level | Size in Memory | Size on Disk | Executors |
|--------------|---------------------------------|----------------|--------------|---------------------------------------------------------------------------------------------------|
| rdd_9_962 | Disk Serialized 2x Replicated | 0.0 B | 388.9 MiB | ip-172-31-3-170.eu-west-1.compute.internal:46697 ip-172-31-3-39.eu-west-1.compute.internal:43111 |
| rdd_9_961 | Disk Serialized 2x Replicated | 0.0 B | 389.1 MiB | ip-172-31-3-196.eu-west-1.compute.internal:39167 ip-172-31-3-36.eu-west-1.compute.internal:46023 |
| rdd_9_960 | Disk Serialized 2x Replicated | 0.0 B | 389.1 MiB | ip-172-31-3-170.eu-west-1.compute.internal:46697 ip-172-31-3-39.eu-west-1.compute.internal:43111 |
| rdd_9_96 | Memory Serialized 2x Replicated | 613.0 MiB | 0.0 B | ip-172-31-3-170.eu-west-1.compute.internal:46697 ip-172-31-3-36.eu-west-1.compute.internal:46023 |
| rdd_9_959 | Disk Serialized 2x Replicated | 0.0 B | 389.5 MiB | ip-172-31-3-170.eu-west-1.compute.internal:46697 ip-172-31-3-36.eu-west-1.compute.internal:46023 |
| rdd_9_958 | Disk Serialized 2x Replicated | 0.0 B | 389.7 MiB | ip-172-31-3-222.eu-west-1.compute.internal:44907 ip-172-31-3-39.eu-west-1.compute.internal:43111 |
| rdd_9_957 | Disk Serialized 2x Replicated | 0.0 B | 388.8 MiB | ip-172-31-3-170.eu-west-1.compute.internal:46697 ip-172-31-3-222.eu-west-1.compute.internal:44907 |
| rdd_9_956 | Disk Serialized 2x Replicated | 0.0 B | 389.9 MiB | ip-172-31-3-222.eu-west-1.compute.internal:44907 ip-172-31-3-39.eu-west-1.compute.internal:43111 |

Environment in Spark Web UI

Environment

▼ Runtime Information

| Name | Value |
|---------------|----------------------------------------------------|
| Java Home | /usr/lib/jvm/java-1.8.0-amazon-corretto.x86_64/jre |
| Java Version | 1.8.0_382 (Amazon.com Inc.) |
| Scala Version | version 2.12.15 |

▼ Spark Properties

| Name | Value |
|-----------------------------------------|------------------------------------------------------------------|
| spark.app.id | application_1694208236041_0001 |
| spark.app.initial.jar.urls | s3://ripani.dub/emr-benchmark-spark/spark-sql-perf.jar |
| spark.app.name | TPCDS Benchmark - c5d.9xlarge - 1GB - maximizeResourceAllocation |
| spark.app.startTime | 1694208354258 |
| spark.app.submitTime | 1694208354217 |
| spark.blacklist.decommissioning.enabled | true |
| spark.blacklist.decommissioning.timeout | 1h |
| spark.decommissioning.timeout.threshold | 20 |
| spark.default.parallelism | 2880 |
| spark.driver.appUIAddress | http://ip-172-31-1-243.eu-west-1.compute.internal:4040 |
| spark.driver.defaultJavaOptions | -XX:OnOutOfMemoryError='kill -9 %p' |

Executors in Spark Web UI

SQL/DataFrame in Spark Web UI

- ❖ SQL/DataFrame page summarizes **all Spark Queries executed in a Spark Application.**
- ❖ This page is only visible in the UI if your application is using Dataset or DataFrame Spark APIs. (Not when using RDD APIs)

APACHE

SPARK

3.4.0-amzn-0

Jobs

Stages

Storage

Environment

Executors

SQL / DataFrame

TPCDS Benchmark - c5d.9xlarge - ... application UI

SQL / DataFrame

Completed Queries: 133

▼ Completed Queries (133)

Page:

<

1

2

3

4

5

6

7

8

9

10

>

>>

14 Pages. Jump to

4

. Show

10

items in a page.

Go

| ID ▼ | Description | Submitted | Duration | Job IDs |
|------|--------------------------------------------------------|---------------------|----------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 102 | benchmark q75-v2.4 <div>+details</div> | 2023/09/08 21:30:26 | 5 s | [724] [725] [726] [727] [728] [729] [730] [731] [732] [733] [734] [735] |
| 101 | benchmark q74-v2.4 <div>+details</div> | 2023/09/08 21:30:23 | 2 s | [713] [714] [715] [716] [717] [718] [719] [720] [721] [722] [723] |
| 100 | benchmark q73-v2.4 <div>+details</div> | 2023/09/08 21:30:21 | 0.6 s | [705] [706] [707] [708] [709] [710] [711] [712] |
| 99 | benchmark q72-v2.4 <div>+details</div> | 2023/09/08 21:30:18 | 2 s | [691] [692] [693] [694] [695] [696] [697] [698] [699] [700] [701] [702] [703] [704] |
| 98 | benchmark q71-v2.4 <div>+details</div> | 2023/09/08 21:30:17 | 0.7 s | [684] [685] [686] [687] [688] [689] [690] |

Details for Query 30

Submitted Time: 2023/09/08 21:27:31

Duration: 1 s

Succeeded Jobs: [90](#) [91](#) [92](#) [93](#) [94](#) [95](#)

☐ Show the Stage ID and Task ID that corresponds to the max metric

Scan parquet

number of files read: 1
scan time: 86 ms
metadata time: 0 ms
size of files read: 1825.0 KiB
max size of file split: 4.0 MiB
number of output rows: 73,049

FileScan parquet

[d_date_sk#492,d_year#498]
Batched: true, DataFilters: [Generate \(2\)](#)
[isnotnull(d_year#498),
(d_year#498 = 2000),
isnotnull(d_date_sk#492)], Format:
Column Parquet, Location:
InMemoryFileIndex(1 paths)
[s3://ripani.dub/warehouse
/tpcds_parquet_1gb/date_dim],
PartitionFilters: [], PushedFilters:
[IsNotNull(d_year),
EqualTo(d_year,2000),
IsNotNull(d_date_sk)],
ReadSchema:
struct<d_date_sk:int,d_year:int>

number of output rows: 366

▼ Details

== Parsed Logical Plan ==

```
GlobalLimit 100
+- 'LocalLimit 100
  +- 'Sort ['i_item_id ASC NULLS FIRST], true
    +- 'Aggregate ['i_item_id], ['i_item_id, 'avg('ss_quantity) AS agg1#27406, 'avg('ss_list_price) AS agg2#27407, 'avg('s
      +- 'Filter (((('ss_sold_date_sk = 'd_date_sk) AND ('ss_item_sk = 'i_item_sk)) AND ('ss_cdemo_sk = 'cd_demo_sk)) AN
        +- 'Join Inner
          :- 'Join Inner
            : :- 'Join Inner
              : : :- 'Join Inner
                : : : :- 'UnresolvedRelation [store_sales], [], false
                  : : : +- 'UnresolvedRelation [customer_demographics], [], false
                    : : +- 'UnresolvedRelation [date_dim], [], false
                      : +- 'UnresolvedRelation [item], [], false
                        +- 'UnresolvedRelation [promotion], [], false
```

== Analyzed Logical Plan ==

```
i_item_id: string, agg1: double, agg2: decimal(11,6), agg3: decimal(11,6), agg4: decimal(11,6)
GlobalLimit 100
+- LocalLimit 100
  +- Sort [i_item_id#565 ASC NULLS FIRST], true
    +- Aggregate [i_item_id#565], [i_item_id#565, avg(ss_quantity#139) AS agg1#27406, avg(ss_list_price#141) AS agg2#27407
      +- Filter (((((ss_sold_date_sk#152 = d_date_sk#492) AND (ss_item_sk#131 = i_item_sk#564)) AND (ss_cdemo_sk#133 = cd
        +- Join Inner
          :- Join Inner
            : :- Join Inner
              : : :- Join Inner
```



```

: : :- Join Inner
: : : :- SubqueryAlias store_sales
: : : : +- View (`store_sales`,
[ss_sold_time_sk#130,ss_item_sk#131,ss_customer_sk#132,ss_cdemo_sk#133,ss_hdemo_sk#134,ss_addr_sk#135,ss_store_sk#136,ss_promo_sk#137]
: : : : +- Relation
[ss_sold_time_sk#130,ss_item_sk#131,ss_customer_sk#132,ss_cdemo_sk#133,ss_hdemo_sk#134,ss_addr_sk#135,ss_store_sk#136,ss_promo_sk#137]
: : : +- SubqueryAlias customer_demographics
: : : : +- View (`customer_demographics`, [cd_demo_sk#474,cd_gender#475,cd_marital_status#476,cd_education_status#477,cd_purchase_mode#478]
: : : : +- Relation [cd_demo_sk#474,cd_gender#475,cd_marital_status#476,cd_education_status#477,cd_purchase_mode#478]
: : +- SubqueryAlias date_dim
: : : +- View (`date_dim`,
[d_date_sk#492,d_date_id#493,d_date#494,d_month_seq#495,d_week_seq#496,d_quarter_seq#497,d_year#498,d_dow#499,d_moy#500,d_dom#501]
: : : +- Relation [d_date_sk#492,d_date_id#493,d_date#494,d_month_seq#495,d_week_seq#496,d_quarter_seq#497,d_year#498,d_dow#499,d_moy#500,d_dom#501]
: +- SubqueryAlias item
: : +- View (`item`, [i_item_sk#564,i_item_id#565,i_rec_start_date#566,i_rec_end_date#567,i_item_desc#568,i_item_category#569,i_item_category_desc#570]
: : : +- Relation [i_item_sk#564,i_item_id#565,i_rec_start_date#566,i_rec_end_date#567,i_item_desc#568,i_item_category#569,i_item_category_desc#570]
+- SubqueryAlias promotion
: : +- View (`promotion`, [p_promo_sk#608,p_promo_id#609,p_start_date_sk#610,p_end_date_sk#611,p_item_sk#612,p_item_category#613,p_item_category_desc#614]
: : : +- Relation [p_promo_sk#608,p_promo_id#609,p_start_date_sk#610,p_end_date_sk#611,p_item_sk#612,p_item_category#613,p_item_category_desc#614]

```

== Optimized Logical Plan ==

GlobalLimit 100

+- LocalLimit 100

+- Sort [i_item_id#565 ASC NULLS FIRST], true

+- Aggregate [i_item_id#565], [i_item_id#565, avg(ss_quantity#139) AS agg1#27406, cast((avg(UnscaledValue(ss_list_price#141) - ss_coupon_amt#148) / ss_sales_price#142) AS DECIMAL(16,4)) AS agg2#27407]

+- Project [ss_quantity#139, ss_list_price#141, ss_sales_price#142, ss_coupon_amt#148, i_item_id#565]

+- Join Inner, (ss_promo_sk#137 = p_promo_sk#608)

:- Project [ss_promo_sk#137, ss_quantity#139, ss_list_price#141, ss_sales_price#142, ss_coupon_amt#148, i_item_id#565]

: +- Join Inner, (ss_item_sk#131 = i_item_sk#564)

Spark Web UI

<https://spark.apache.org/docs/latest/web-ui.html>

AWS EMR Best Practices

<https://aws.github.io/aws-emr-best-practices/docs/bestpractices/>